Advanced Topics in Social and Geographic Data Science

GPS data classification

UCL

GPS data classification

- classification (supervised learning problem involving predicting a class label)
- generative model (using a decision boundary)
- tree-based methods (boosted decision tree, random forest)
- support vector machines

Decision surface

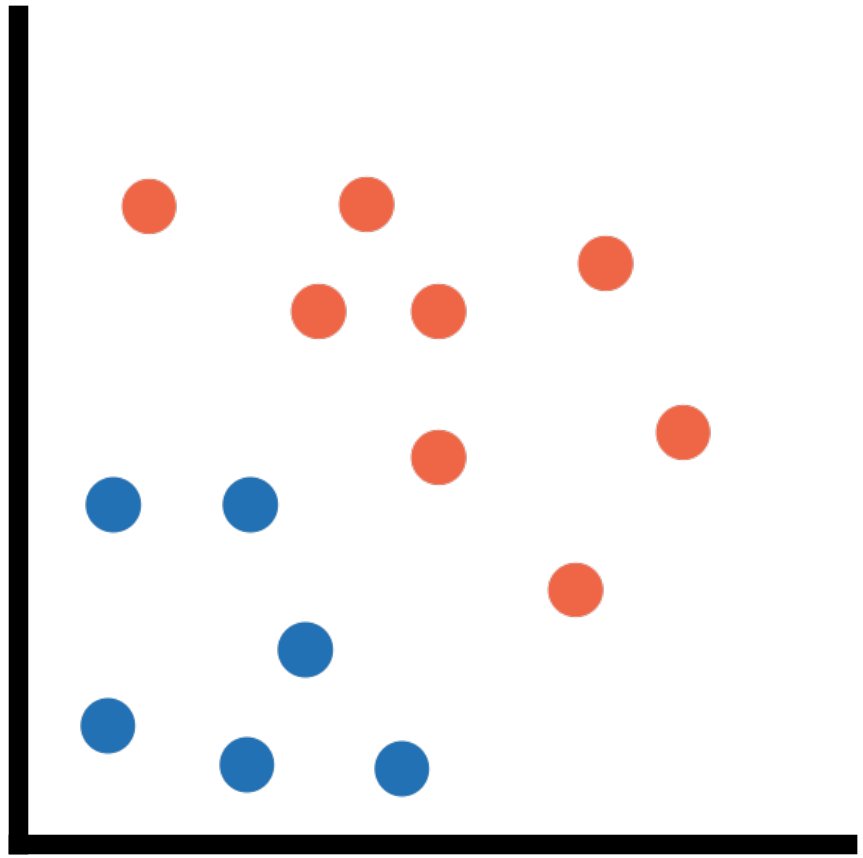

Decision surface

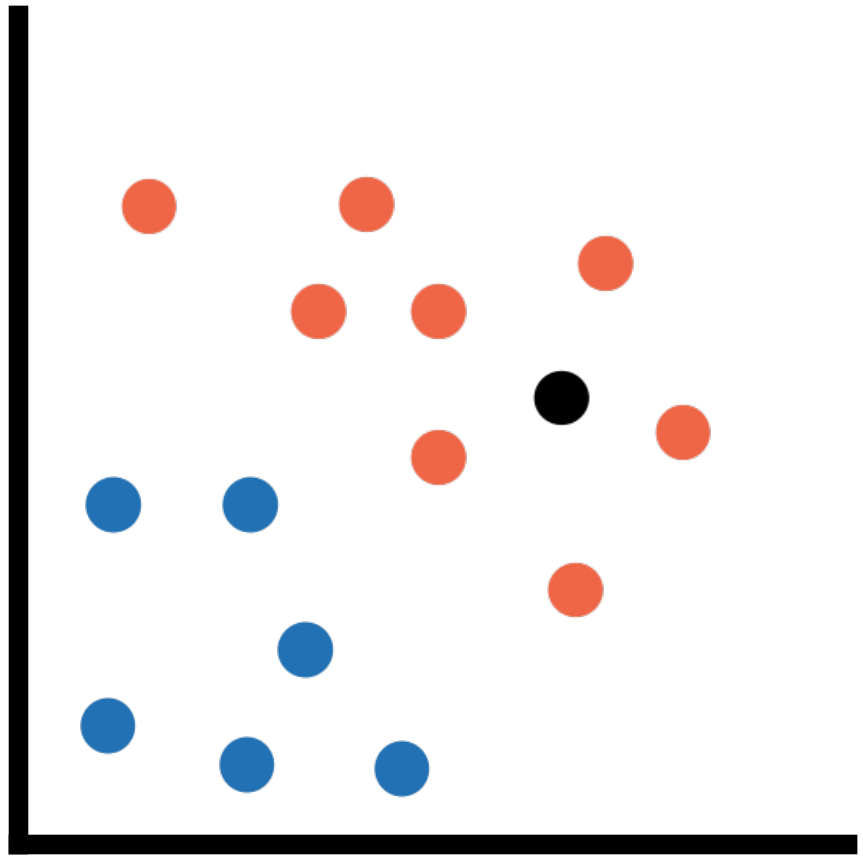

Decision surface

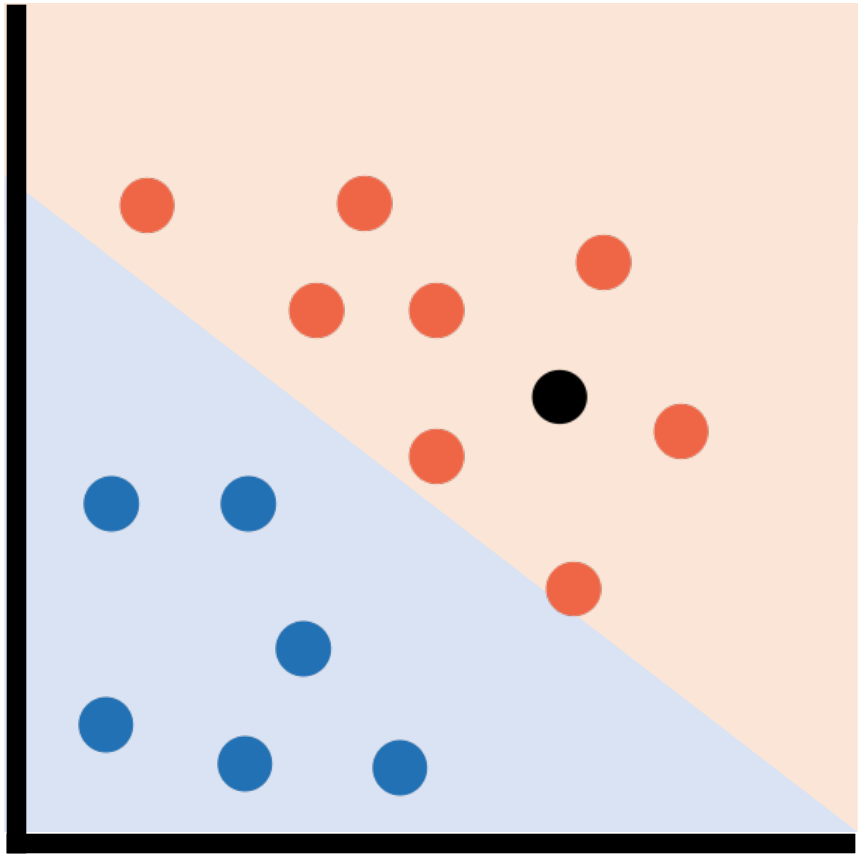

Processing pipeline

- split into train and test
- develop model on the train data set (decision surface)
- predict classes on the test data set
- measure of accuracy: percentage correctly predicted or Kappa value for unbalanced classes (comparison of prediction to class probabilities)

Decision trees

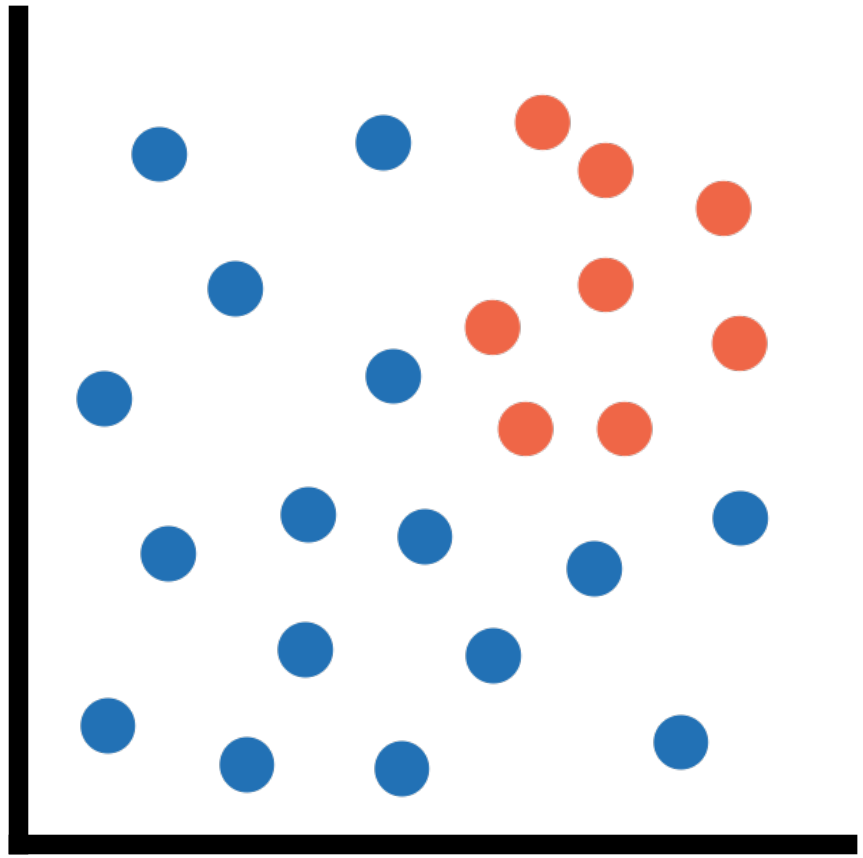

Decision trees

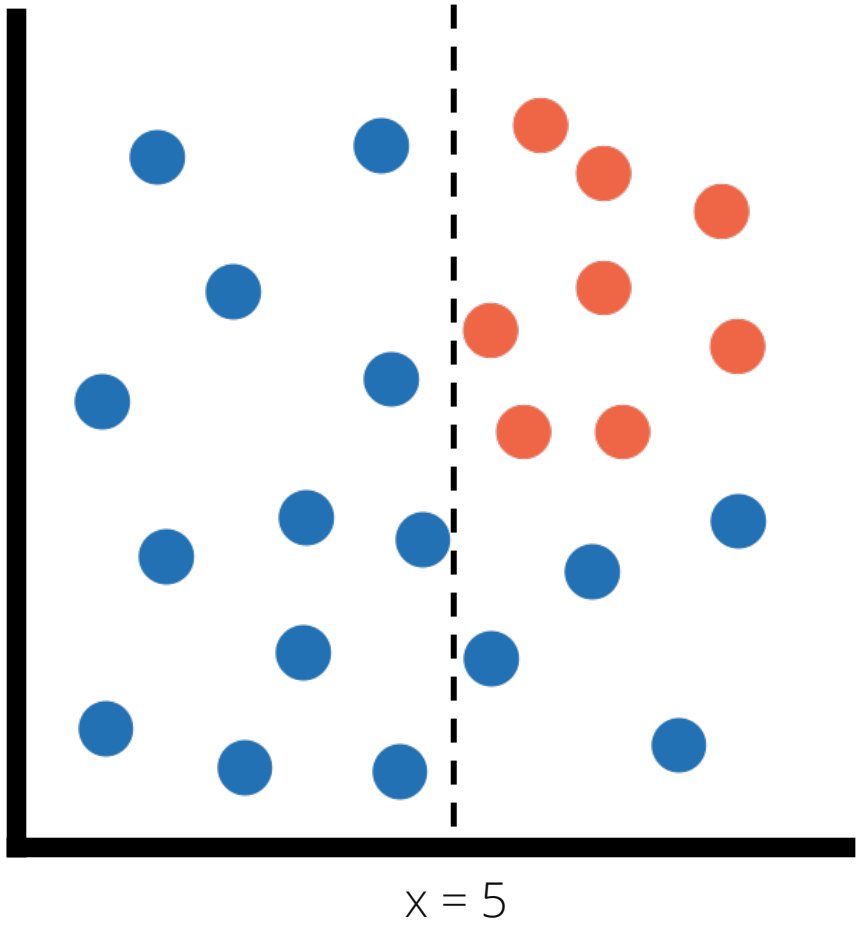

Decision trees

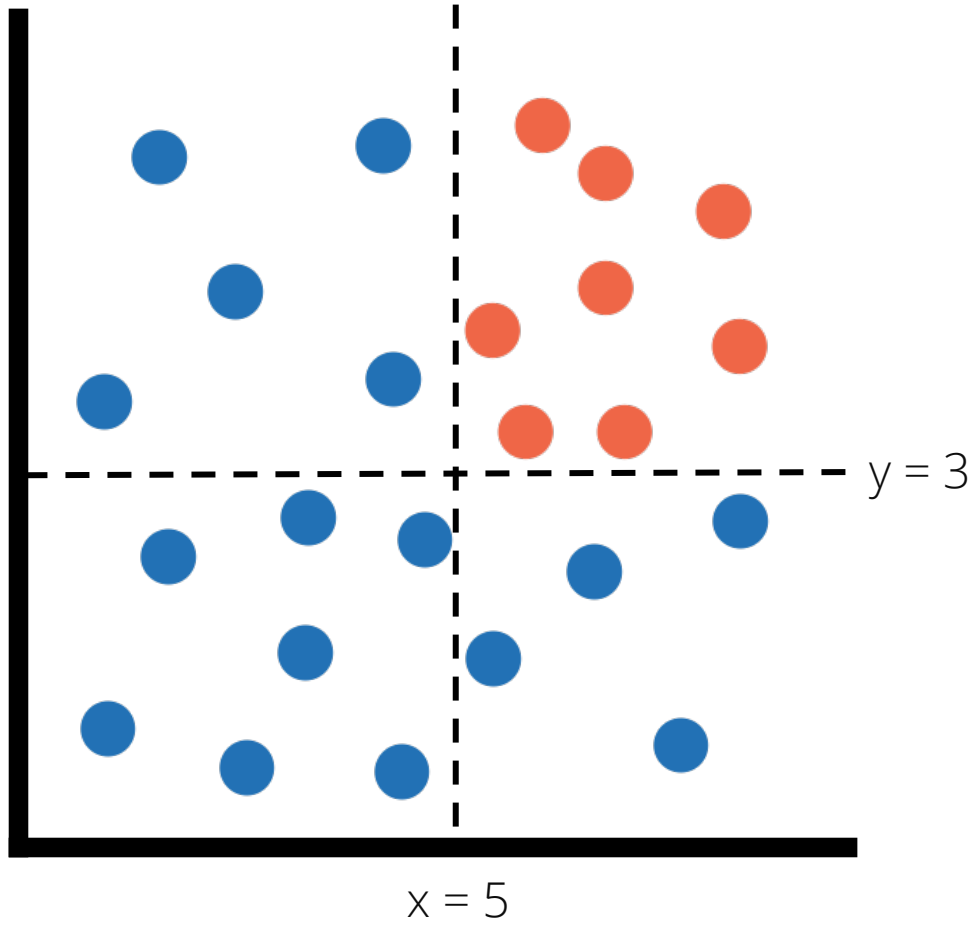

Decision trees

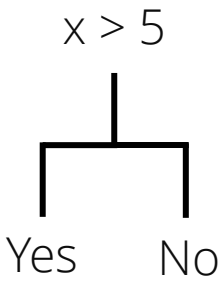

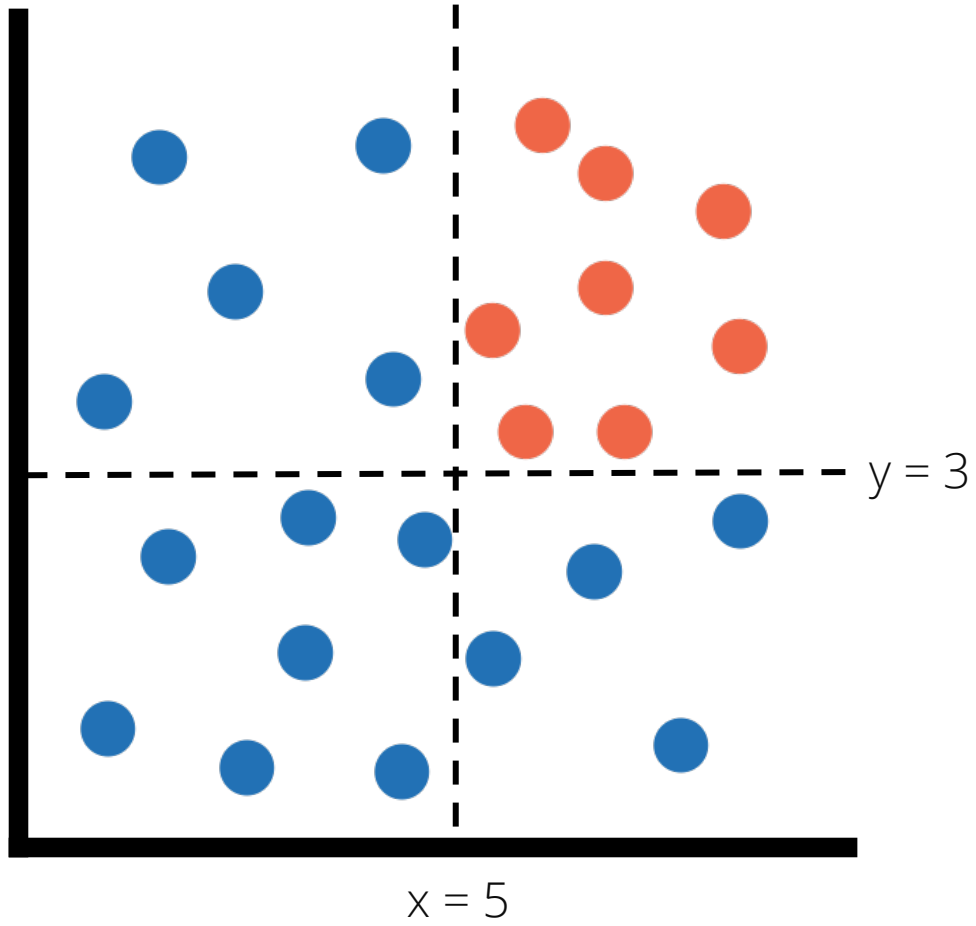

Decision trees

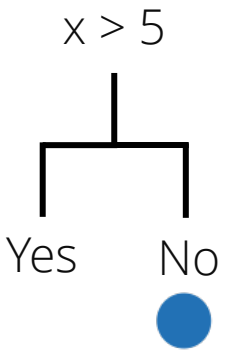

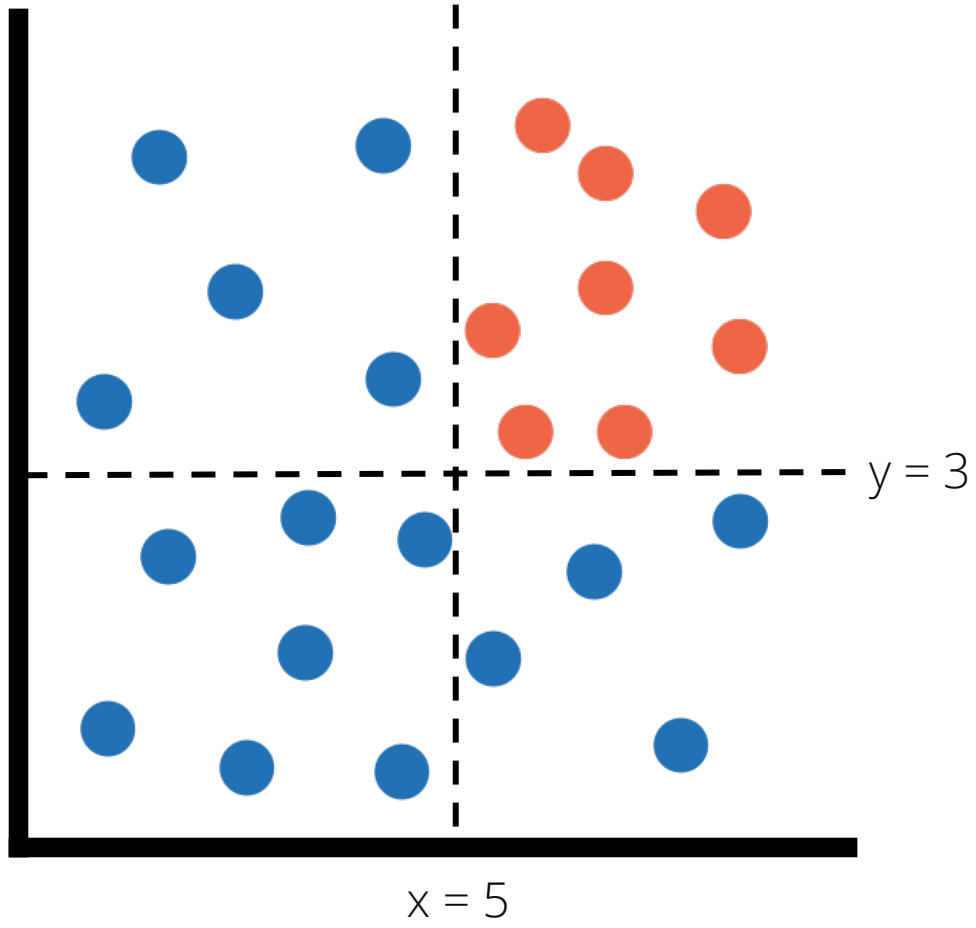

Decision trees

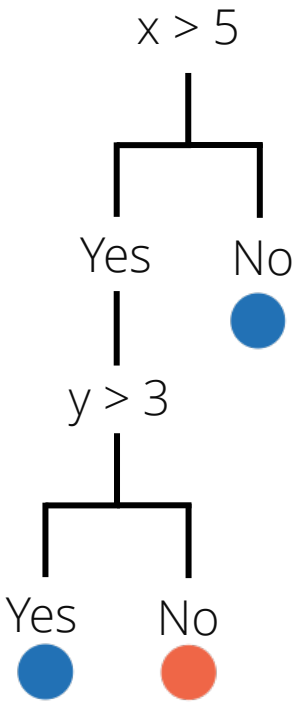

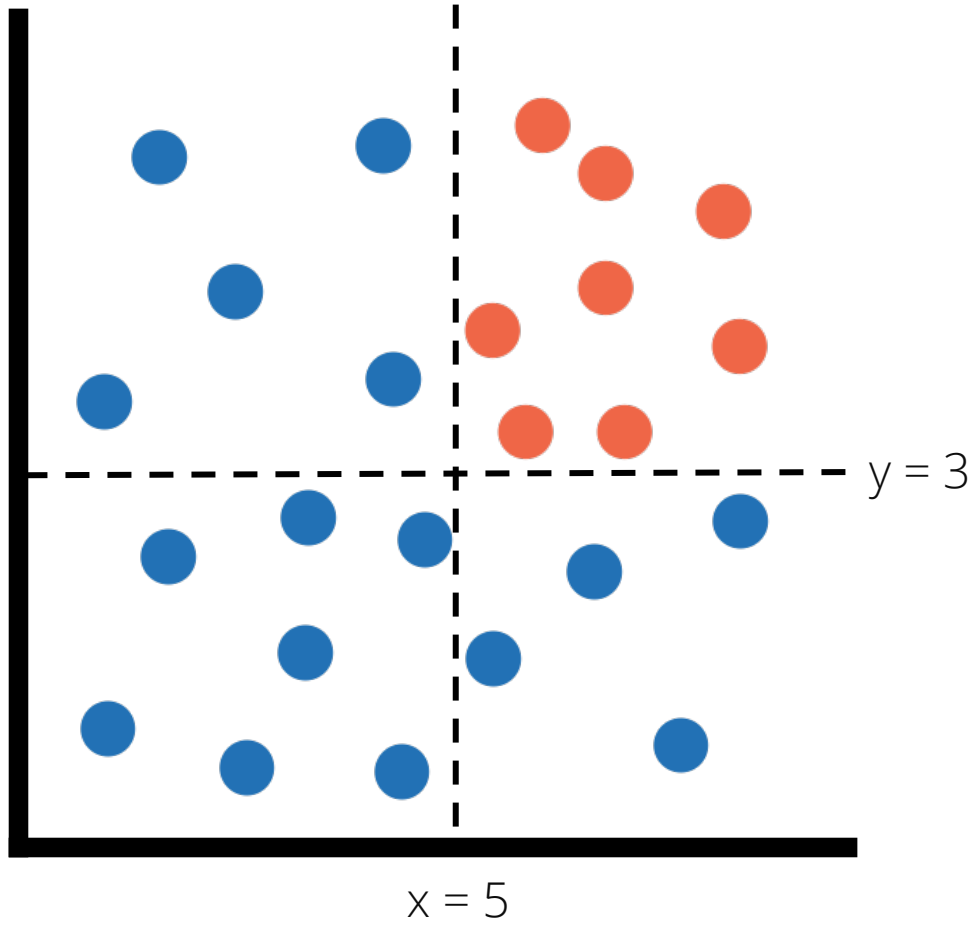

Decision trees

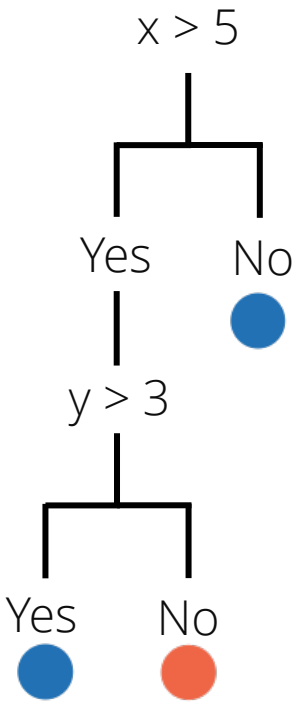

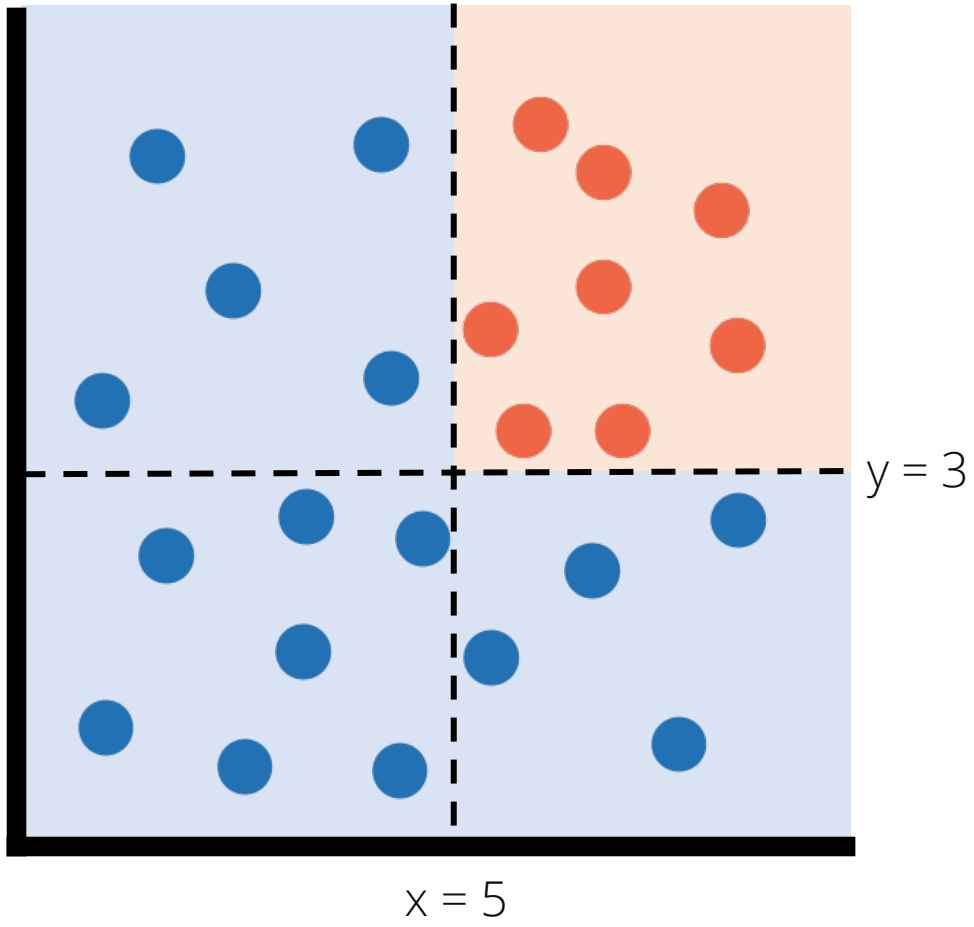

Splitting the data

- entropy used to split the data (how to decide where to split the data): measure of impurity (which split will give the best results), i.e. is there a split possible where one of branches gets all class label of one group
  - all examples in same class: entropy is 0
  - all examples equally split: entropy is 1
- entropy used to calculate information gain – which split will lead to the largest decrease in entropy

Typical parameters

- minimum sample split (when to stop splitting the data to avoid overfitting): minimum number of samples in your test data at the leaves of your tree
- pruning the tree (removing sections with little predictive power) to avoid overfitting

Boosted decision trees

- grow a full decision tree
- test the resulting tree against the training data
- try to improve the misclassified class labels with a new decision tree
- test the resulting tree against the training data
- try to improve the misclassified class labels with a new decision tree
- …and repeat till maximum number of boosting iterations has been reached

- for prediction: majority vote for final class label

Random forest

- grow many decisions trees (an ensemble of trees) but on a subset of the training data using bootstrap sampling (random sampling with replacement)
- at every split an *n* number of randomly selected variables is used to calculate the next split
- result is a forest of trees that have been grown using varying features on different subsets of the training data

- for prediction: majority vote for final class labels

Tree-based methods

- many different implementations, different parameters
- different ways of calculating information gain
- typically involve a form of boosting, bagging

Support Vector Machines

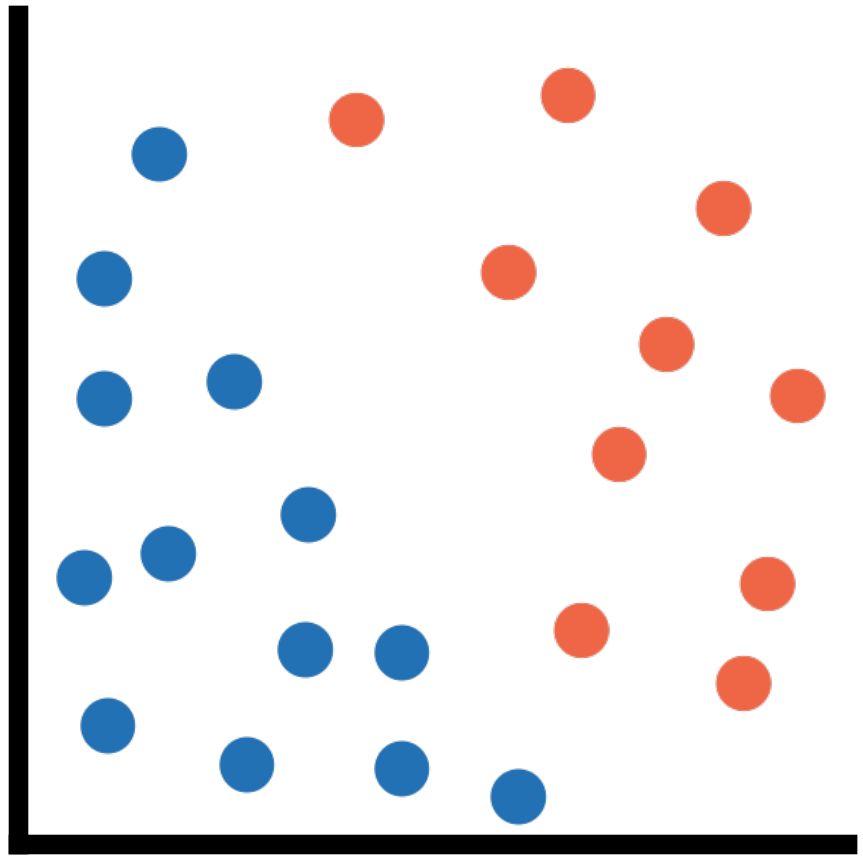

Support Vector Machines

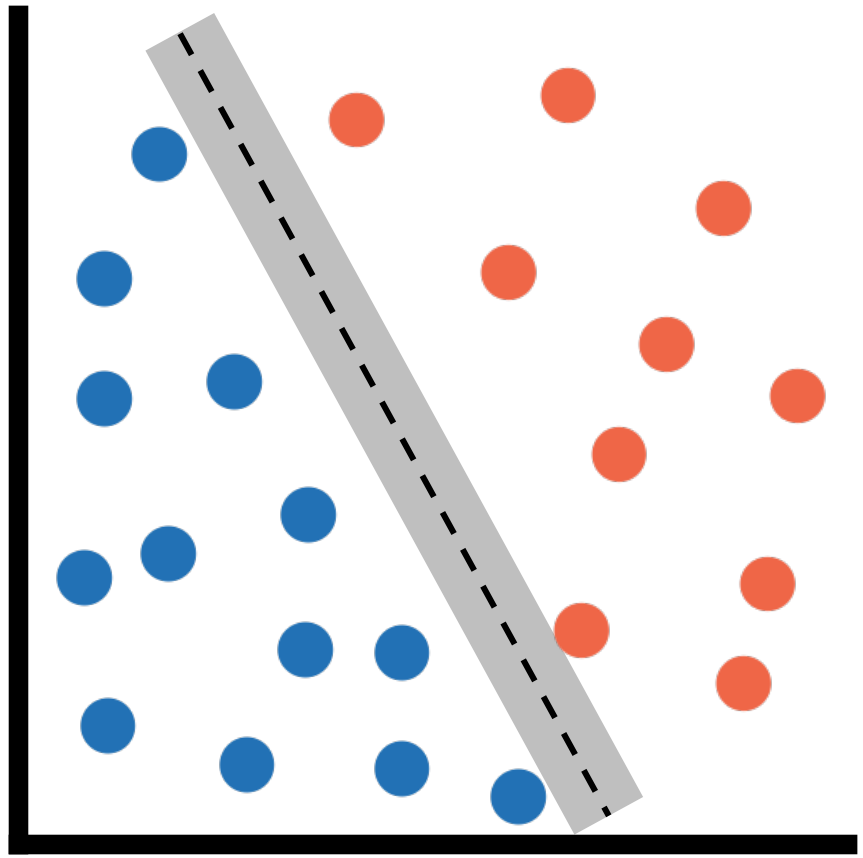

Support Vector Machines

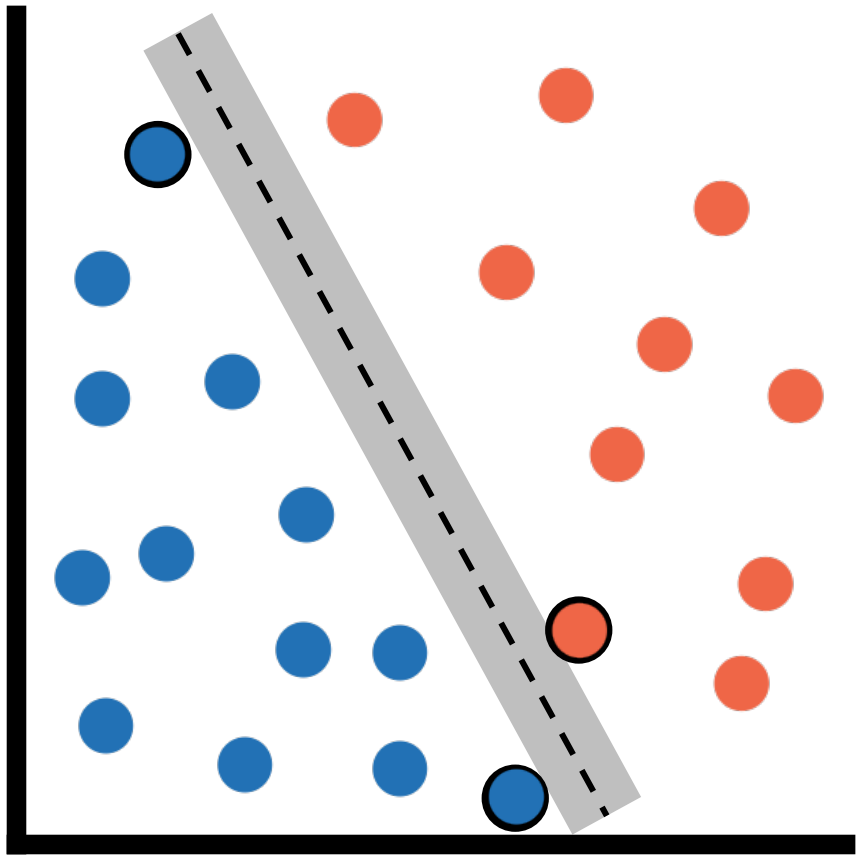

Support Vector Machines

- tries to find the largest possible hyperplane to linearly separate the data by maximising the distance to the nearest points (margin)
- different way to deal with outliers (ignoring)
- SVM is a classifier built on linear-separation but it can generate non-linear decision boundaries by transforming the input data to make them linear separatable through something called the kernel trick
- kernel trick transforms low-dimensional data into high-dimensional data to make them linearly separatable

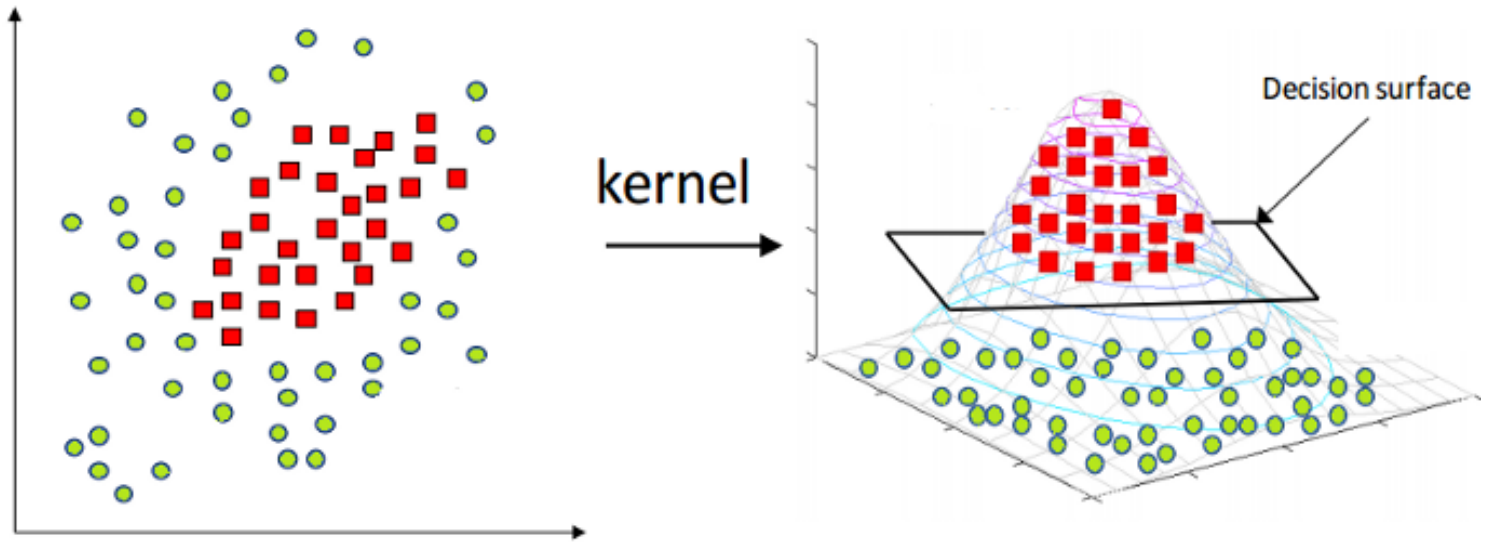
kernel
Decision surface

Support Vector Machines

- different kernels
- C parameter (controls the tradeoff between smooth decision boundary and one that classifies all the training points correctly)
- Gamma parameter (defines the range of influence of a single training example has on the decision boundaries)

GPS data classification

- brief introduction to two sets of supervised machine learning classifiers that could be used to impute trips (moves) and activities (stays) from raw GPS trajectories
- like all supervised machine learning classifiers: labelled data are required to train the classifiers